

```
In [171... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [172... import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)
```

```
In [173... %matplotlib inline
```

```
In [174... df = pd.read_csv('HW2.csv')
df.head()
```

/var/folders/qs/5vt90tvd159ctgq252mwn07r0000gn/T/ipykernel_92460/3949404760.py:1: DtypeWarning: Columns (26) have mixed types. Specify dtype option on import or set low_memory=False.

```
df = pd.read_csv('HW2.csv')
```

```
Out[174]:
```

	ID	Customer_ID	Month	Name	Age	SSN	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank_Accounts
--	----	-------------	-------	------	-----	-----	------------	---------------	-----------------------	-------------------

0	0x1602	CUS_0xd40	January	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	1824.843333	3
1	0x1603	CUS_0xd40	February	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	NaN	3
2	0x1604	CUS_0xd40	March	Aaron Maashoh	-500	821-00-0265	Scientist	19114.12	NaN	3
3	0x1605	CUS_0xd40	April	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	NaN	3
4	0x1606	CUS_0xd40	May	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	1824.843333	3

5 rows x 28 columns

```
In [175... df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 28 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                     100000 non-null object
1   Customer_ID                           100000 non-null object
2   Month                                  100000 non-null object
3   Name                                   90015 non-null  object
4   Age                                    100000 non-null object
5   SSN                                    100000 non-null object
6   Occupation                             100000 non-null object
7   Annual_Income                           100000 non-null object
8   Monthly_Inhand_Salary                   84998 non-null float64
9   Num_Bank_Accounts                       100000 non-null int64
10  Num_Credit_Card                          100000 non-null int64
11  Interest_Rate                           100000 non-null int64
12  Num_of_Loan                              100000 non-null object
13  Type_of_Loan                             88592 non-null object
14  Delay_from_due_date                      100000 non-null int64
15  Num_of_Delayed_Payment                   92998 non-null object
16  Changed_Credit_Limit                     100000 non-null object
17  Num_Credit_Inquiries                     98035 non-null float64
18  Credit_Mix                               100000 non-null object
19  Outstanding_Debt                         100000 non-null object
20  Credit_Utilization_Ratio                 100000 non-null float64
21  Credit_History_Age                       90970 non-null object
22  Payment_of_Min_Amount                    100000 non-null object
23  Total_EMI_per_month                      100000 non-null float64
24  Amount_invested_monthly                  95521 non-null object
25  Payment_Behaviour                        100000 non-null object
26  Monthly_Balance                          98800 non-null object
27  Credit_Score                             100000 non-null object
dtypes: float64(4), int64(4), object(20)
memory usage: 21.4+ MB
```

```
In [176... df2 = df[[column for column in df if df[column].count()/len(df) >= 0.3]]
del df2['ID']
print("List of dropped columns:", end=" ")
for c in df.columns:
    if c not in df2.columns:
        print(c, end = ', ')
print('\n')
df = df2
```

List of dropped columns: ID,

```
In [177... print(df.isnull().values.any())  
print(df.shape)
```

```
True  
(1000000, 27)
```

```
In [178... print(df.isnull().sum())
```

```
Customer_ID          0  
Month                0  
Name                9985  
Age                 0  
SSN                 0  
Occupation           0  
Annual_Income        0  
Monthly_Inhand_Salary 15002  
Num_Bank_Accounts    0  
Num_Credit_Card      0  
Interest_Rate        0  
Num_of_Loan          0  
Type_of_Loan         11408  
Delay_from_due_date  0  
Num_of_Delayed_Payment 7002  
Changed_Credit_Limit 0  
Num_Credit_Inquiries 1965  
Credit_Mix          0  
Outstanding_Debt     0  
Credit_Utilization_Ratio 0  
Credit_History_Age   9030  
Payment_of_Min_Amount 0  
Total_EMI_per_month  0  
Amount_invested_monthly 4479  
Payment_Behaviour    0  
Monthly_Balance      1200  
Credit_Score         0  
dtype: int64
```

```
In [179... df.dropna(inplace=True)
```

Since the frequency of missing values are not too significant, we can decide to drop the null values and still be able to find some insights.

```
In [180... for col in df:  
    print(df[col].describe())  
    print("\n")
```

```
count          53049
unique          11071
top            CUS_0x533b
freq              8
Name: Customer_ID, dtype: object
```

```
count          53049
unique              8
top            January
freq            6744
Name: Month, dtype: object
```

```
count          53049
unique          9168
top            Stevex
freq             30
Name: Name, dtype: object
```

```
count          53049
unique          1057
top              38
freq            1567
Name: Age, dtype: object
```

```
count          53049
unique          11064
top            #F%$D@*&8
freq            2971
Name: SSN, dtype: object
```

```
count          53049
unique           16
top
freq            3727
Name: Occupation, dtype: object
```

```
count          53049
unique          14707
top            20867.67
```

```
freq          11  
Name: Annual_Income, dtype: object
```

```
count    53049.000000  
mean      4024.502617  
std       3094.967739  
min        303.645417  
25%       1575.345833  
50%       2991.016667  
75%       5712.553333  
max      15204.633333  
Name: Monthly_Inhand_Salary, dtype: float64
```

```
count    53049.000000  
mean       16.842806  
std       115.605025  
min        -1.000000  
25%         4.000000  
50%         6.000000  
75%         8.000000  
max      1798.000000  
Name: Num_Bank_Accounts, dtype: float64
```

```
count    53049.000000  
mean       22.787253  
std       129.339640  
min         0.000000  
25%         4.000000  
50%         6.000000  
75%         7.000000  
max      1499.000000  
Name: Num_Credit_Card, dtype: float64
```

```
count    53049.000000  
mean       76.272352  
std       479.226709  
min         1.000000  
25%         8.000000  
50%        15.000000  
75%        22.000000  
max      5797.000000
```

Name: Interest_Rate, dtype: float64

```
count    53049
unique     249
top         3
freq     8644
```

Name: Num_of_Loan, dtype: object

```
count    53049
unique    6259
top    Not Specified
freq      874
```

Name: Type_of_Loan, dtype: object

```
count    53049.000000
mean      21.994119
std       15.220406
min       -5.000000
25%       10.000000
50%       19.000000
75%       29.000000
max       67.000000
```

Name: Delay_from_due_date, dtype: float64

```
count    53049
unique     470
top        19
freq     3169
```

Name: Num_of_Delayed_Payment, dtype: object

```
count    53049
unique    3975
top       _
freq     1102
```

Name: Changed_Credit_Limit, dtype: object

```
count    53049.000000
mean      27.062866
std      187.185825
```

```
min          0.000000
25%          3.000000
50%          6.000000
75%          9.000000
max         2594.000000
Name: Num_Credit_Inquiries, dtype: float64
```

```
count        53049
unique         4
top      Standard
freq         19418
Name: Credit_Mix, dtype: object
```

```
count        53049
unique       11407
top        1151.7
freq         16
Name: Outstanding_Debt, dtype: object
```

```
count        53049.000000
mean         32.219031
std          5.057398
min         20.881250
25%         28.046338
50%         32.251786
75%         36.408991
max         49.564519
Name: Credit_Utilization_Ratio, dtype: float64
```

```
count          53049
unique          404
top    18 Years and 2 Months
freq          257
Name: Credit_History_Age, dtype: object
```

```
count        53049
unique         3
top          Yes
freq        29670
Name: Payment_of_Min_Amount, dtype: object
```



```

count      53049.000000
mean       1445.924761
std        8407.429893
min         4.462837
25%        41.244418
50%        78.418272
75%        169.286269
max        82331.000000
Name: Total_EMI_per_month, dtype: float64

```

```

count      53049
unique     50540
top        __10000__
freq       2402
Name: Amount_invested_monthly, dtype: object

```

```

count      53049
unique      7
top        Low_spent_Small_value_payments
freq       13581
Name: Payment_Behaviour, dtype: object

```

```

count      53049
unique     53044
top        __-333333333333333333333333333333__
freq       6
Name: Monthly_Balance, dtype: object

```

```

count      53049
unique      3
top        Standard
freq       27951
Name: Credit_Score, dtype: object

```

```

In [181... df = df[~df['Occupation'].str.contains("_")]
df = df[df['Num_Bank_Accounts']>0]
df = df[df['Delay_from_due_date']>0]

```

```

df = df[~df['Age'].astype(str).str.contains('_')]
df = df[df['Age'].astype(int)>0]
df = df[df['Age'].astype(int)<100]
df = df[~df['Amount_invested_monthly'].astype(str).str.contains('_')]
df = df[~df['Monthly_Balance'].astype(str).str.contains('_')]
df = df[~df['Changed_Credit_Limit'].astype(str).str.contains('_')]
df = df[~df['SSN'].astype(str).str.contains('#F%')]
df = df[~df['Num_of_Loan'].astype(str).str.contains('_')]
df = df[df['Num_of_Loan'].astype(int)>0]
df = df[~df['Credit_Mix'].astype(str).str.contains('_')]
df = df[~df['Payment_Behaviour'].astype(str).str.contains('!@')]
df = df[df['Changed_Credit_Limit'].astype(float)>0]
df = df[df['Num_Credit_Inquiries'].astype(int)<40]
df = df[df['Num_Credit_Inquiries'].astype(int)>0]
df = df[df['Num_of_Loan'].astype(int)<40]
df = df[df['Num_Bank_Accounts']<20]
df = df[df['Num_Credit_Card']<20]
df = df[df['Num_Credit_Card']>0]
df = df[df['Interest_Rate']<100]
df = df[df['Total_EMI_per_month']<2000.0]
df = df[~df['Num_of_Delayed_Payment'].astype(str).str.contains('_')]
df = df[df['Num_of_Delayed_Payment'].astype(int)<100]
df = df[~df['Annual_Income'].astype(str).str.contains('_')]
df = df[~df['Outstanding_Debt'].astype(str).str.contains('_')]
df = df[df['Annual_Income'].astype(float)<600000.0]

```

```

In [182... df['Changed_Credit_Limit'] = df['Changed_Credit_Limit'].astype(float).round(3)
df['Total_EMI_per_month'] = df['Total_EMI_per_month'].round(3)
df['Credit_Utilization_Ratio'] = df['Credit_Utilization_Ratio'].round(3)
df['Amount_invested_monthly'] = df['Amount_invested_monthly'].astype(float).round(3)
df['Monthly_Balance'] = df['Monthly_Balance'].astype(float).round(3)
df['Monthly_Inhand_Salary'] = df['Monthly_Inhand_Salary'].round(3)
df['Annual_Income'] = df['Annual_Income'].astype(float).round(3)

```

```

In [183... '''
Now let's change the data type from object to int/float for the columns:
int--
Age
Num_of_Loan
Delay_from_due_date

float--
Outstanding_Debt
'''

```

```
df['Age'] = df['Age'].astype('int64')
df['Num_of_Loan'] = df['Num_of_Loan'].astype('int64')
df['Delay_from_due_date'] = df['Delay_from_due_date'].astype('int64')
df['Num_of_Delayed_Payment'] = df['Num_of_Delayed_Payment'].astype('int64')
df['Outstanding_Debt'] = df['Outstanding_Debt'].astype('float64')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 17431 entries, 9 to 99991
```

```
Data columns (total 27 columns):
```

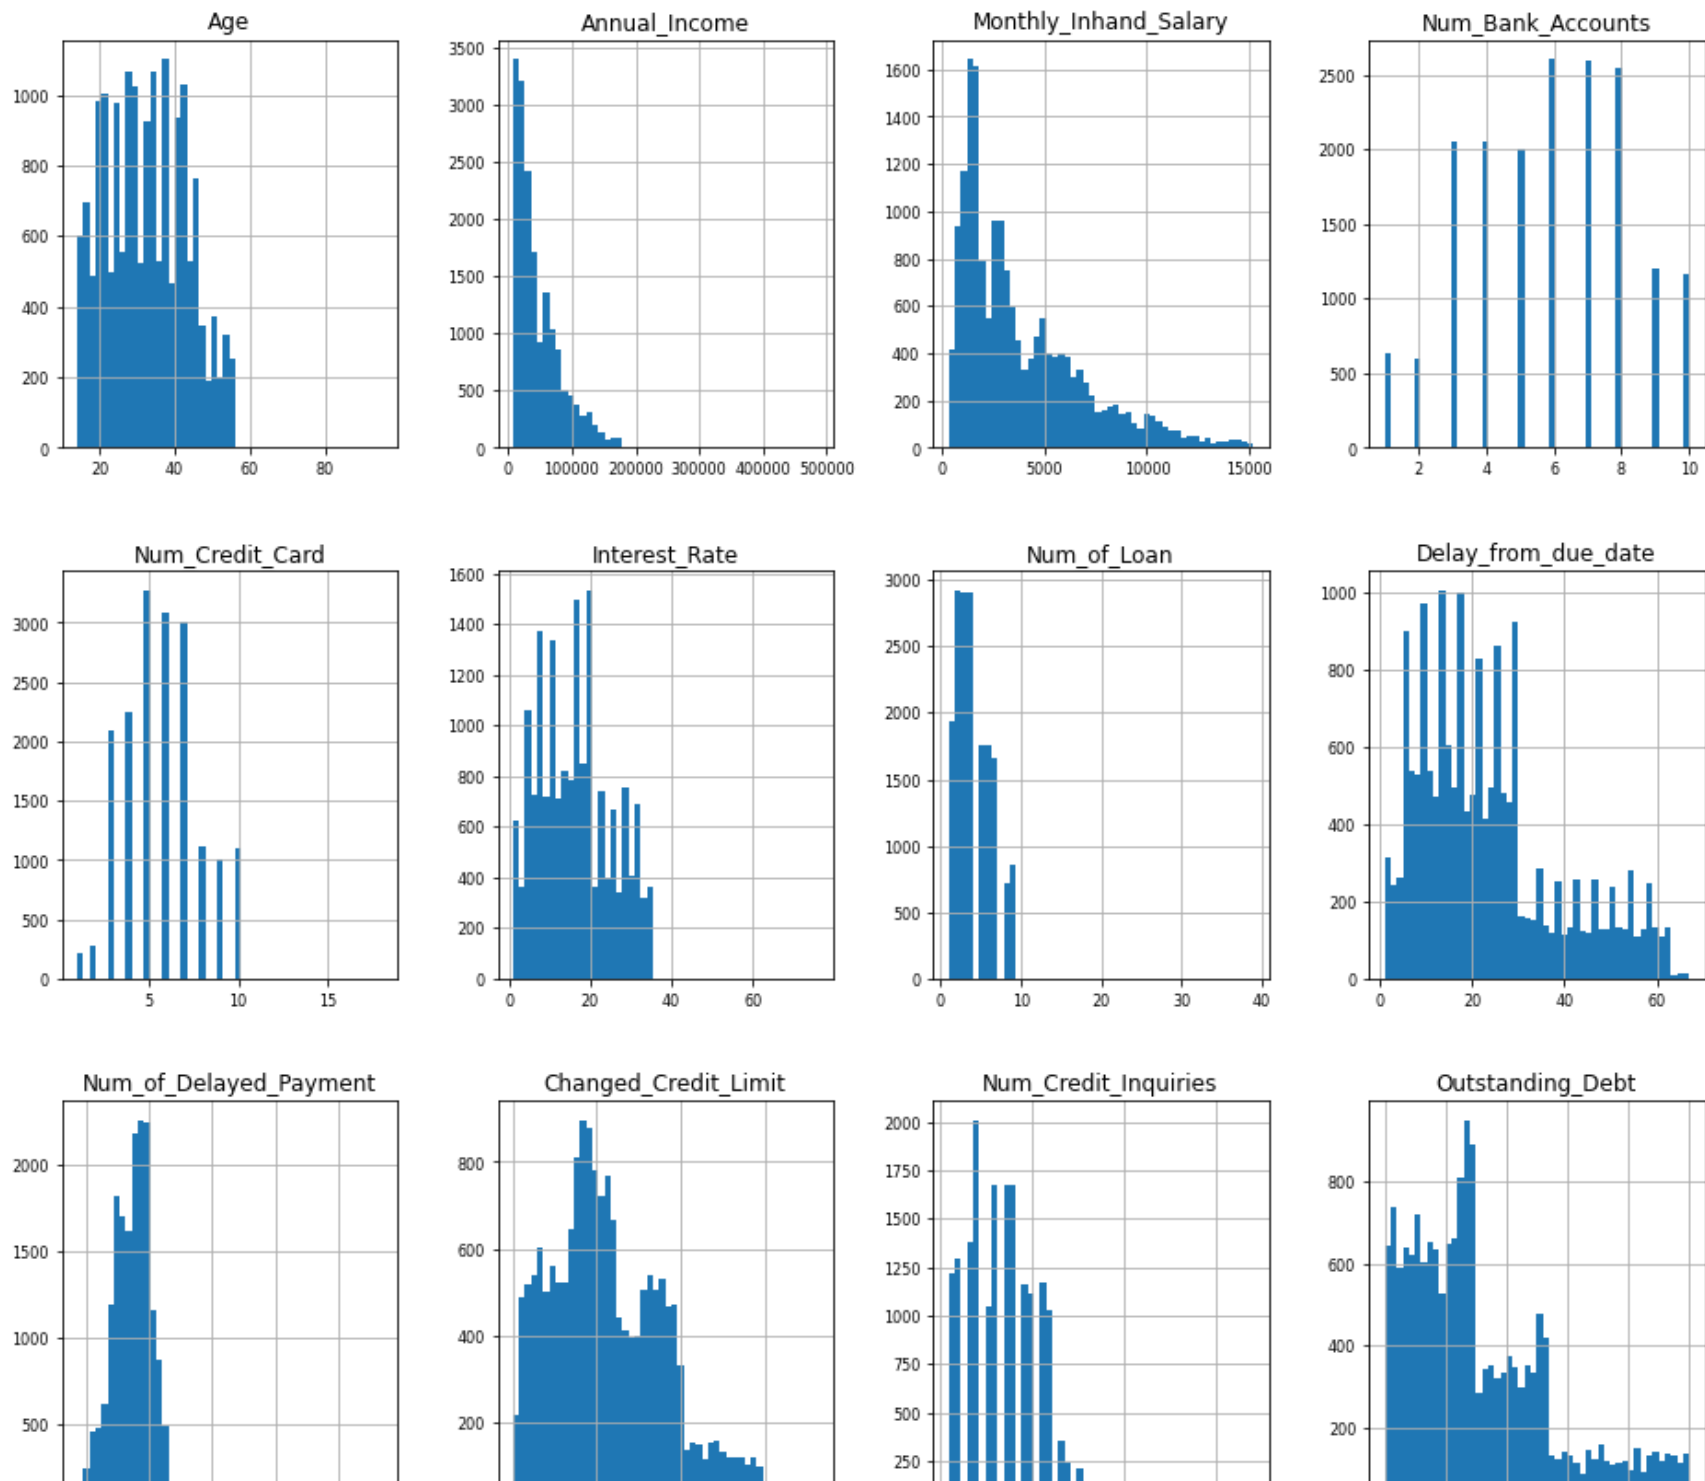
#	Column	Non-Null Count	Dtype
0	Customer_ID	17431 non-null	object
1	Month	17431 non-null	object
2	Name	17431 non-null	object
3	Age	17431 non-null	int64
4	SSN	17431 non-null	object
5	Occupation	17431 non-null	object
6	Annual_Income	17431 non-null	float64
7	Monthly_Inhand_Salary	17431 non-null	float64
8	Num_Bank_Accounts	17431 non-null	int64
9	Num_Credit_Card	17431 non-null	int64
10	Interest_Rate	17431 non-null	int64
11	Num_of_Loan	17431 non-null	int64
12	Type_of_Loan	17431 non-null	object
13	Delay_from_due_date	17431 non-null	int64
14	Num_of_Delayed_Payment	17431 non-null	int64
15	Changed_Credit_Limit	17431 non-null	float64
16	Num_Credit_Inquiries	17431 non-null	float64
17	Credit_Mix	17431 non-null	object
18	Outstanding_Debt	17431 non-null	float64
19	Credit_Utilization_Ratio	17431 non-null	float64
20	Credit_History_Age	17431 non-null	object
21	Payment_of_Min_Amount	17431 non-null	object
22	Total_EMI_per_month	17431 non-null	float64
23	Amount_invested_monthly	17431 non-null	float64
24	Payment_Behaviour	17431 non-null	object
25	Monthly_Balance	17431 non-null	float64
26	Credit_Score	17431 non-null	object

```
dtypes: float64(9), int64(7), object(11)
```


```
memory usage: 3.7+ MB
```

```
In [184.. df_num = df.select_dtypes(include=['int64', 'float64'])
df_num.hist(figsize=(16,20), bins=50, xlabelsize=8, ylabelsize=8)
```

```
Out[184]: array([[<AxesSubplot:title={'center': 'Age'}>,
                  <AxesSubplot:title={'center': 'Annual_Income'}>,
                  <AxesSubplot:title={'center': 'Monthly_Inhand_Salary'}>,
                  <AxesSubplot:title={'center': 'Num_Bank_Accounts'}>],
                [<AxesSubplot:title={'center': 'Num_Credit_Card'}>,
                  <AxesSubplot:title={'center': 'Interest_Rate'}>,
                  <AxesSubplot:title={'center': 'Num_of_Loan'}>,
                  <AxesSubplot:title={'center': 'Delay_from_due_date'}>],
                [<AxesSubplot:title={'center': 'Num_of_Delayed_Payment'}>,
                  <AxesSubplot:title={'center': 'Changed_Credit_Limit'}>,
                  <AxesSubplot:title={'center': 'Num_Credit_Inquiries'}>,
                  <AxesSubplot:title={'center': 'Outstanding_Debt'}>],
                [<AxesSubplot:title={'center': 'Credit_Utilization_Ratio'}>,
                  <AxesSubplot:title={'center': 'Total_EMI_per_month'}>,
                  <AxesSubplot:title={'center': 'Amount_invested_monthly'}>,
                  <AxesSubplot:title={'center': 'Monthly_Balance'}>]], dtype=object)
```



```
In [185... corr_matrix = df_num.corr()
plt.figure(figsize=(15,15))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', square=True)
plt.title("Correlation Matrix of Numerical Features")
plt.show()
```



Given the strong correlation between "Annual_Income" and "Monthly_Inhand_Salary", it is advisable to drop one of the variables. In this case, I will remove "Annual_Income" as the distribution of "Monthly_Inhand_Salary" appears to be more balanced.

```
In [186... df.drop(columns=['Annual_Income'], inplace=True)  
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 17431 entries, 9 to 99991
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Customer_ID                          17431 non-null  object
1   Month                                17431 non-null  object
2   Name                                  17431 non-null  object
3   Age                                  17431 non-null  int64
4   SSN                                   17431 non-null  object
5   Occupation                           17431 non-null  object
6   Monthly_Inhand_Salary                17431 non-null  float64
7   Num_Bank_Accounts                    17431 non-null  int64
8   Num_Credit_Card                      17431 non-null  int64
9   Interest_Rate                        17431 non-null  int64
10  Num_of_Loan                          17431 non-null  int64
11  Type_of_Loan                         17431 non-null  object
12  Delay_from_due_date                  17431 non-null  int64
13  Num_of_Delayed_Payment               17431 non-null  int64
14  Changed_Credit_Limit                 17431 non-null  float64
15  Num_Credit_Inquiries                 17431 non-null  float64
16  Credit_Mix                           17431 non-null  object
17  Outstanding_Debt                    17431 non-null  float64
18  Credit_Utilization_Ratio             17431 non-null  float64
19  Credit_History_Age                   17431 non-null  object
20  Payment_of_Min_Amount                17431 non-null  object
21  Total_EMI_per_month                  17431 non-null  float64
22  Amount_invested_monthly              17431 non-null  float64
23  Payment_Behaviour                    17431 non-null  object
24  Monthly_Balance                      17431 non-null  float64
25  Credit_Score                         17431 non-null  object
dtypes: float64(8), int64(7), object(11)
memory usage: 3.6+ MB
```

```
In [187... df.shape
```

```
Out[187]: (17431, 26)
```

```
In [189... ohe_features = ['Customer_ID',
                  'Month',
                  'Name',
                  'SSN',
                  'Occupation'
                  ]
```

```
te_features = [ 'Type_of_Loan',
                 'Credit_Mix',
                 'Credit_History_Age',
                 'Payment_of_Min_Amount',
                 'Payment_Behaviour',
                 ]

num_features = [ 'Age',
                 'Monthly_Balance',
                 'Amount_invested_monthly',
                 'Outstanding_Debt',
                 'Changed_Credit_Limit',
                 'Num_of_Delayed_Payment',
                 'Monthly_Inhand_Salary',
                 'Num_Bank_Accounts',
                 'Num_Credit_Card',
                 'Interest_Rate',
                 'Num_of_Loan',
                 'Delay_from_due_date',
                 'Num_Credit_Inquiries',
                 'Credit_Utilization_Ratio',
                 'Total_EMI_per_month' ]

feature_names = ohe_features + te_features + num_features
```

```
In [188... # Encode non-numeric values
categorical_cols = ['Credit_Score'] + ohe_features + te_features
for col in categorical_cols:
    if df[col].dtype == 'object':
        df[col] = df[col].astype('category').cat.codes
```

```

Customer_ID      0
Month            0
Name             0
Age             0
SSN             0
Occupation       0
Monthly_Inhand_Salary  0
Num_Bank_Accounts  0
Num_Credit_Card  0
Interest_Rate    0
Num_of_Loan      0
Type_of_Loan     0
Delay_from_due_date  0
Num_of_Delayed_Payment  0
Changed_Credit_Limit  0
Num_Credit_Inquiries  0
Credit_Mix       0
Outstanding_Debt  0
Credit_Utilization_Ratio  0
Credit_History_Age  0
Payment_of_Min_Amount  0
Total_EMI_per_month  0
Amount_invested_monthly  0
Payment_Behaviour  0
Monthly_Balance  0
Credit_Score     0
dtype: int64

```

```

In [196... from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder
from sklearn.compose import make_column_transformer
from category_encoders import TargetEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import make_pipeline

X_dev, X_test, y_dev, y_test = train_test_split(df[feature_names], df[['Credit_Score']], stratify=df[['Credit_
X_train, X_val, y_train, y_val = train_test_split(X_dev, y_dev, stratify=y_dev, test_size=0.2, random_state=4

lg = LogisticRegression(multi_class='multinomial', solver='lbfgs', max_iter=1000)
preprocess = make_column_transformer((StandardScaler(), num_features),

```

```
        (OneHotEncoder(handle_unknown='ignore'), ohe_features),
        (TargetEncoder(handle_unknown='ignore'), te_features),
        remainder='passthrough')
pipe = make_pipeline(preprocess, GridSearchCV(lg,
                                              param_grid={},
                                              return_train_score=True))

pipe.fit(X_train, y_train)
grid_search_results = pipe.named_steps['gridsearchcv']
print(f"Best train score: ", grid_search_results.best_score_)
print(f"Best train alpha: ", grid_search_results.best_params_)
print(f"Test score: ", pipe.score(X_dev, y_dev))
```

Warning: No categorical columns found. Calling 'transform' will only return input data.

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
    y = column_or_1d(y, warn=True)
```

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
    n_iter_i = _check_optimize_result(
```

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
    y = column_or_1d(y, warn=True)
```

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
    n_iter_i = _check_optimize_result(
```

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
    y = column_or_1d(y, warn=True)
```

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
    n_iter_i = _check_optimize_result(
```

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
    y = column_or_1d(y, warn=True)
```

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/validation.py:993: DataConversionWarn
ing: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceW
arning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/validation.py:993: DataConversionWarn
ing: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

Best train score: 0.6649036306588973

Best train alpha: {}

Test score: 0.6645152036718301

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceW
arning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

During the model development phase, a validation dataset is used to improve the performance of a model by evaluating its performance on a subset of the original data that is held out during training. This dataset is used to fine-tune the model's hyperparameters and select the best performing model. While the validation dataset is helpful in improving the model's performance during training, it should not be used to obtain a final estimate of the model's performance on new, unseen data. Instead, a separate test dataset should be used for this purpose to ensure an unbiased estimate of the model's performance.

```
In [192... y_pred = pipe.predict(X_val)
val_acc = accuracy_score(y_val, y_pred)
val_report = classification_report(y_val, y_pred)
print("Validation accuracy:", val_acc)
print("Validation classification report:", val_report)
```

Validation accuracy: 0.6457511652922194

Validation classification report:				precision	recall	f1-score	support
0	0.56	0.46	0.51	373			
1	0.67	0.54	0.60	945			
2	0.65	0.76	0.70	1471			
accuracy				0.65			2789
macro avg				0.63	0.59	0.60	2789
weighted avg				0.65	0.65	0.64	2789

```
In [193... y_pred = pipe.predict(X_test)
test_acc = accuracy_score(y_test, y_pred)
test_report = classification_report(y_test, y_pred)
print("Test accuracy:", test_acc)
print("Test classification report:", test_report)
```

Test accuracy: 0.6498422712933754

Test classification report:				precision	recall	f1-score	support
0	0.56	0.46	0.51	467			
1	0.67	0.55	0.61	1181			
2	0.66	0.76	0.70	1839			
accuracy				0.65			3487
macro avg				0.63	0.59	0.61	3487
weighted avg				0.65	0.65	0.64	3487

Based on the classification report, the model has an overall test accuracy of 0.65, meaning it correctly predicted 65% of the samples in the test set.

The precision for class 0 is 0.56, meaning that out of all the samples the model predicted to be in class 0, only 56% of them are actually in class 0. The recall for class 0 is 0.46, meaning that out of all the samples in class 0, the model correctly predicted 46% of them. The F1-score for class 0 is 0.51, which is the harmonic mean of precision and recall.

Similarly, for class 1, the precision is 0.67, the recall is 0.55, and the F1-score is 0.61. For class 2, the precision is 0.66, the recall is 0.76, and the F1-score is 0.70.

The macro-average F1-score is 0.61, which is the average of the F1-scores for all three classes. The weighted-average F1-score is 0.64, which takes into account the imbalance in the number of samples in each class. Overall, the model performs reasonably well but may need further improvements, especially for predicting class 0.

CART has several benefits, such as producing easily interpretable decision trees, processing both numerical and categorical data for classification and regression, and handling missing data through exclusion during tree construction. However, the algorithm is susceptible to overfitting when the tree grows too deep or when the stopping criterion is not optimized.

```
In [204... from sklearn.tree import DecisionTreeClassifier, plot_tree

dtc = DecisionTreeClassifier(random_state=82)
preprocess = make_column_transformer((StandardScaler(), num_features),
                                     (OneHotEncoder(handle_unknown='ignore'), ohe_features),
                                     (TargetEncoder(handle_unknown='ignore'), te_features),
                                     remainder='passthrough')
pipe = make_pipeline(preprocess, GridSearchCV(dtc,
                                             param_grid={},
                                             return_train_score=True))

pipe.fit(X_train, y_train)
grid_search_results = pipe.named_steps['gridsearchcv']
print(f"Best train score: ", grid_search_results.best_score_)
print(f"Best train alpha: ", grid_search_results.best_params_)
print(f"Test score:", pipe.score(X_dev, y_dev))
```

```
Warning: No categorical columns found. Calling 'transform' will only return input data.
Best train score: 0.6893769610040341
Best train alpha: {}
Test score: 0.9386115892139989
```

```
In [205... best_tree = grid_search_results.best_estimator_
ohe_feature_names = preprocess.named_transformers_['onehotencoder'].get_feature_names()
te_feature_names = preprocess.named_transformers_['targetencoder'].get_feature_names()
feature_names = num_features + ohe_feature_names.tolist() + te_feature_names
visual_tree = plot_tree(best_tree, feature_names=feature_names, filled=True)
plt.show()
```


/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function get_feature_names is deprecated; get_feature_names is deprecated in 1.0 and will be removed in 1.2. Please use get_feature_names_out instead.

warnings.warn(msg, category=FutureWarning)

/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/category_encoders/utils.py:360: FutureWarning: `get_feature_names` is deprecated in all of sklearn. Use `get_feature_names_out` instead.

warnings.warn("`get_feature_names` is deprecated in all of sklearn. Use `get_feature_names_out` instead.",



```
In [206... test_accuracies = []

for i in range(30):
    dtc = DecisionTreeClassifier(random_state=i)
    preprocess = make_column_transformer((StandardScaler(), num_features),
                                         (OneHotEncoder(handle_unknown='ignore'), ohe_features),
                                         (TargetEncoder(handle_unknown='ignore'), te_features),
                                         remainder='passthrough')
    pipe = make_pipeline(preprocess, GridSearchCV(dtc,
                                                  param_grid={},
                                                  return_train_score=True))

    pipe.fit(X_train, y_train)
    test_accuracy = pipe.score(X_dev, y_dev)
    test_accuracies.append(test_accuracy)
    print(f"Run {i}: test score = {test_accuracy}")

test_accuracies = np.array(test_accuracies)
mean_accuracy = np.mean(test_accuracies)
std_accuracy = np.std(test_accuracies)

print(f"Average accuracy: {mean_accuracy}")
print(f"Standard deviation: {std_accuracy}")
```

Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 0: test score = 0.9403327596098681
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 1: test score = 0.9395438898450946
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 2: test score = 0.9389701663798049
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 3: test score = 0.9389701663798049
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 4: test score = 0.9398307515777395
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 5: test score = 0.9386833046471601
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 6: test score = 0.9386833046471601
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 7: test score = 0.9397590361445783
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 8: test score = 0.9400458978772231
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 9: test score = 0.9388984509466437
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 10: test score = 0.9399024670109007
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 11: test score = 0.9394004589787722
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 12: test score = 0.9396873207114171
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 13: test score = 0.9395438898450946
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 14: test score = 0.9398307515777395
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 15: test score = 0.9403327596098681
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 16: test score = 0.9391135972461274
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 17: test score = 0.9381095811818704
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 18: test score = 0.9391853126792886
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 19: test score = 0.9399024670109007
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 20: test score = 0.9390418818129661
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 21: test score = 0.9399024670109007
Warning: No categorical columns found. Calling 'transform' will only return input data.

```

Run 22: test score = 0.9392570281124498
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 23: test score = 0.9397590361445783
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 24: test score = 0.9386833046471601
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 25: test score = 0.9388984509466437
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 26: test score = 0.9392570281124498
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 27: test score = 0.9405479059093517
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 28: test score = 0.939328743545611
Warning: No categorical columns found. Calling 'transform' will only return input data.
Run 29: test score = 0.9401893287435457
Average accuracy: 0.9394530502964239
Standard deviation: 0.0005745993187061979

```

Both bagging and random forests are powerful ensemble learning techniques that utilize multiple decision trees to enhance model performance. Bagging generates independent bootstrap samples of the training data, and fits a decision tree to each sample, with predictions being aggregated across all trees in the ensemble. On the other hand, random forests use a specific type of bagging, where each decision tree is trained on a random subset of features, reducing the risk of overfitting and increasing generalization performance. Random forests also employ feature bagging, or random subspace method, to add more randomness to the model, and further reduce overfitting. Overall, random forests are a more sophisticated approach than bagging for boosting model performance.

Feature subsampling, a key technique in random forest models, trains each decision tree on a random subset of features, reducing correlation between trees, avoiding overfitting, and increasing model performance. By enabling the model to focus on the most important features while ignoring the noise, this technique is efficient in handling high-dimensional and complex data.

```

In [203... from sklearn.ensemble import RandomForestClassifier

rfc = RandomForestClassifier(random_state=0)
param_grid = {}
pipe = make_pipeline(preprocess, GridSearchCV(rfc,
                                              param_grid = param_grid,
                                              return_train_score=True))

pipe.fit(X_train, y_train)
grid_search_results = pipe.named_steps['gridsearchcv']
print(f"Best train score: ", grid_search_results.best_score_)

```

```
print(f"Best train alpha: ", grid_search_results.best_params_)
print(f"Test score:", pipe.score(X_dev, y_dev))
```

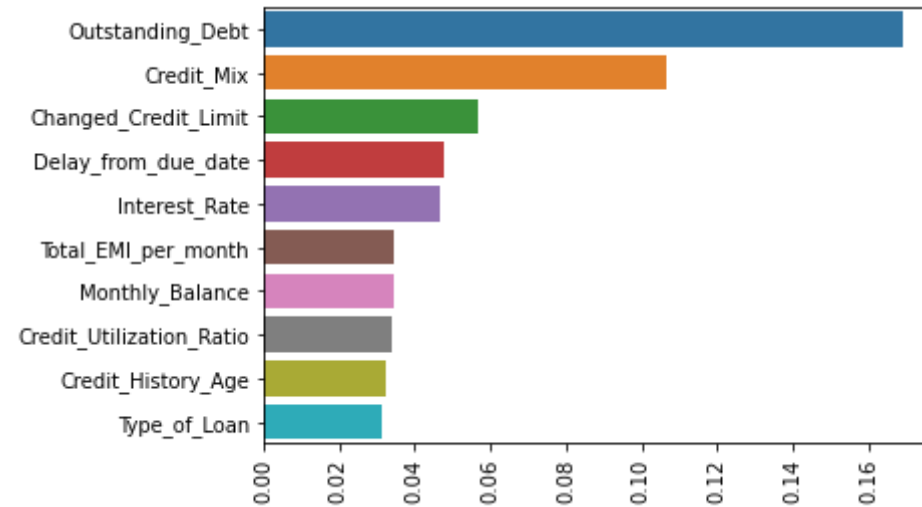
Warning: No categorical columns found. Calling 'transform' will only return input data.

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_validation.py:680: DataCon
versionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n
_samples,), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_validation.py:680: DataCon
versionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n
_samples,), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_validation.py:680: DataCon
versionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n
_samples,), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_validation.py:680: DataCon
versionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n
_samples,), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_validation.py:680: DataCon
versionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n
_samples,), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_search.py:926: DataConvers
ionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_sam
ples,), for example using ravel().
    self.best_estimator_.fit(X, y, **fit_params)
Best train score: 0.7367099955177051
Best train alpha: {}
Test score: 0.9491537578886976
```

Random forest had a higher train/test score than those of decision tree classifiers and logistic regression models. The results of this evaluation indicate that the random forest model outperformed the other two models in terms of its train/test score. This suggests that the random forest algorithm was better able to capture the underlying patterns and relationships within the data and make more accurate predictions than the decision tree classifiers and logistic regression models.

```
In [207... rf = grid_search_results.best_estimator_
featimps = zip(feature_names, rf.feature_importances_)
feats, imps = zip(*(sorted(list(filter(lambda x: x[1] != 0, featimps)), key=lambda x: x[1], reverse=True)))
ax = sns.barplot(list(imps[:10]), list(feats[:10]))
ax.tick_params(axis='x', rotation=90)
```

```
/Users/jasonjin/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(
```



In []:

In []: